

TRAINING COURSE

Apache Kafka® Administration by Confluent

Course Objectives

During this hands-on course you will learn how:

- Kafka and the Confluent platform work, and how their main subsystems interact
- To set up, manage, monitor and tune your cluster
- To use industry best practices developed by the world's foremost Apache Kafka® experts

Hands-on Training

Throughout the course, hands-on exercises reinforce the topics being discussed.

Exercises include:

- Using Kafka's command-line tools
- Automating configuration
- Using Kafka's administrative tools
- Tuning Producer and Consumer performance
- Securing the cluster
- Building data pipelines with Kafka Connect

Prerequisites

Attendees should have a working knowledge of the Kafka architecture, either through:

- Prior experience, or
- By taking Confluent Fundamentals for Apache Kafka®, which can be accessed [here](#).

It is also important to have strong knowledge of Linux/Unix and understand basic TCP/IP networking concepts. Familiarity with Java Virtual Machine (JVM) is helpful.

Participants are required to provide a laptop computer with unobstructed internet access to fully participate in the class.

To evaluate your Kafka knowledge for this course, please complete the self assessment:

<https://cnfl.io/fundamentals-quiz>

To sign-up for one of our courses, visit us [here](#).

Who Should Attend?

This course is designed for engineers, system administrators, and operations staff responsible for building, managing, monitoring, and tuning Kafka clusters.

Content

MODULE	DESCRIPTION
Fundamentals of Apache Kafka®	<ul style="list-style-type: none"> • Kafka as a Distributed Streaming Platform • The Distributed Log • Producer and Consumer Basics
Apache Kafka® Architecture	<ul style="list-style-type: none"> • Kafka's Commit Log • Replication for High Availability • Partitions and Consumer Groups for Scalability • Security Overview
Providing Durability	<ul style="list-style-type: none"> • Data Replication • Failure Recovery • Log Files & Offset Management • Exactly-Once Semantics (EOS)
Managing a Cluster	<ul style="list-style-type: none"> • Installing and Running Kafka • Configuration Management • Monitoring • Log Retention and Compaction • Commissioning and Decommissioning Brokers
Optimizing Apache Kafka® Performance	<ul style="list-style-type: none"> • Monitoring, Testing, and Tuning Brokers and Kafka Clients • The Consumer Group Protocol
Apache Kafka® Security	<ul style="list-style-type: none"> • Transport Encryption • Authentication • Securing Apache Kafka® • Migrating to a Secure Cluster
Data Pipelines with Kafka Connect	<ul style="list-style-type: none"> • The Motivation for Kafka Connect • Types of Connectors • Kafka Connect Implementation • Standalone and Distributed Modes • Configuring the Connectors
Kafka in Production	<ul style="list-style-type: none"> • Kafka Reference Architecture for Apache Kafka® and the Complete Confluent Platform • Capacity Planning • Multi Data Center Deployments

Confluent offers instructor-led courses in both traditional and virtual classroom formats, as well as in a self-paced format available through the Confluent Self-Paced Subscription. Visit confluent.io/training for more information.

Disclaimer: Subscriptions purchased with Training Credits will have an end date matching the Training Credits order.